

Deep Learning System and Parallel Computing HW1

- Author: 黃柏瑄 (P78081528)

1. Experiment description

- Framework: Chainer
- Dataset: Cifar10
- Model: ResNet20

Train ResNet20 with Cifar10 dataset on Chainer without GPU and test the inference result.

2. Environment setting

- Hardware:
 - CPU: Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz (`$ lscpu | grep 'Model name'`)
 - Memory: 32 GB
- Software:
 - OS: Ubuntu 18.04.5 LTS (`$ lsb_release -d`)
 - Python 3.6.9
 - Chainer 7.7.0

```
$ virtualenv --python python3.6 env
$ source env/bin/activate
$ pip install chainer==7.7.0
$ pip list
Package           Version
-----
chainer            7.7.0
filelock           3.0.12
numpy              1.19.2
pip                20.2.2
protobuf           3.13.0
setuptools         50.2.0
six                1.15.0
typing-extensions 3.7.4.3
wheel              0.35.1
```

3. Result

Code can be found at [Appendix](#) section.

3.1. Train

```
$ python train.py
```

```
[env] .../benlab/pro/ms/dpc2020/hw1/ [main ↑1 •1 010 x] [20:33]
$ python train.py
GPU: -1
# unit: 10
# Minibatch-size: 128
# epoch: 10

epoch      main/loss  validation/main/loss  main/accuracy  validation/main/accuracy  elapsed_time
1          1.49377    1.21208              0.447011       0.563884                 1022.99
2          1.067     1.24985              0.616828       0.561511                 2049.73
3          0.887509   0.890478             0.686759       0.691456                 3092.96
4          0.757574   0.960231             0.733156       0.665941                 4115.04
5          0.662063   0.770418             0.767104       0.728936                 5068.98
6          0.587038   0.83906              0.795533       0.715783                 5988.97
7          0.531226   0.729016             0.815018       0.754252                 6926.88
8          0.486883   0.706037             0.829788       0.766515                 7837.67
9          0.434821   0.65807              0.848326       0.778283                 8782.04
10         0.399522   0.700709             0.860394       0.773339                 9703.1

[env] .../benlab/pro/ms/dpc2020/hw1/ [main ↑1 •1 07 x] [23:16]
```

3.2. Inference

```
$ python inference.py
```

```
$ python inference.py
# unit: 10

The test data label: [3]
result: variable([[ -2.0492222 -3.4841754 -1.5581383 -1.2748083 -2.167873
                  -1.9728453 -2.962061 -2.8587072 -3.0967042 -2.4422274]])
result Top 1: [3]
result Top 5: [3 2 5 0 4]
```

4. Comprehension and discovery

4.1. ResNet Model

For different dataset, the model will be a little different. Take ResNet as the example, ResNet20, 32, 44, 56, 110 is used for Cifar10 whose BottleNeck has 2 or 3 Convolution layers. On the other hand, ResNet50, 101, 152, which perform on ImageNet, has 3 or 4 Convolution per BottleNeck block.

4.2. CPU utilization

The CPU utilization is good during training phase.

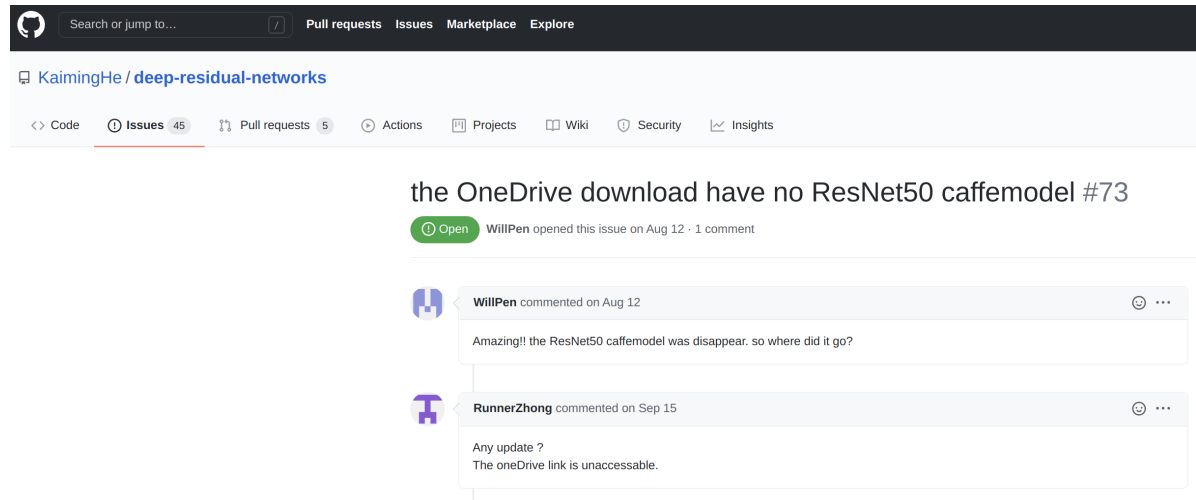
```
1  [|||||] 100.0% 5  [|||||] 100.0% 9  [|||||] 100.0% 13  [|||||] 100.0%
2  [|||||] 100.0% 6  [|||||] 100.0% 10 [|||||] 100.0% 14  [|||||] 100.0%
3  [|||||] 100.0% 7  [|||||] 100.0% 11 [|||||] 100.0% 15  [|||||] 100.0%
4  [|||||] 100.0% 8  [|||||] 100.0% 12 [|||||] 100.0% 16  [|||||] 100.0%
Mem[|||||] 18.2G/31.3G Tasks: 261, 1673 thr: 16 running
Swp[|||||] 754M/30.5G Load average: 11.39 8.58 4.37
Uptime: 19 days, 21:42:52

PID PRI VIRT RES SHR S CPU% MEM% TIME+ Command
14846 20 5949M 3905M 16680 R 1570 12.2 51:41.52 python train.py
15036 20 5949M 3905M 16680 R 99.1 12.2 2:57.07 python train.py
15032 20 5949M 3905M 16680 R 98.5 12.2 2:57.26 python train.py
15038 20 5949M 3905M 16680 R 98.5 12.2 2:55.99 python train.py
15028 20 5949M 3905M 16680 R 98.5 12.2 2:57.53 python train.py
15030 20 5949M 3905M 16680 R 98.5 12.2 2:57.35 python train.py
15025 20 5949M 3905M 16680 R 98.5 12.2 2:58.77 python train.py
15034 20 5949M 3905M 16680 R 98.5 12.2 2:56.84 python train.py
15033 20 5949M 3905M 16680 R 97.8 12.2 2:57.91 python train.py
15035 20 5949M 3905M 16680 R 97.8 12.2 2:56.92 python train.py
15027 20 5949M 3905M 16680 R 97.8 12.2 2:59.26 python train.py
15031 20 5949M 3905M 16680 R 97.8 12.2 2:56.96 python train.py
15037 20 5949M 3905M 16680 R 97.8 12.2 2:56.78 python train.py
15029 20 5949M 3905M 16680 R 97.2 12.2 2:57.57 python train.py
15039 20 5949M 3905M 16680 R 97.2 12.2 2:54.45 python train.py
15026 20 5949M 3905M 16680 R 95.9 12.2 2:58.77 python train.py
```

4.3. Fun facts

4.3.1. It is not easy to find some pre-trained weights on the Internet

Many pre-trained models are disappear. There may be many reasons such as privacy issue, business considerations, or lack of maintenance.



source: <https://github.com/KaimingHe/deep-residual-networks/issues/73>

4.3.2. Preferred Networks Migrates its Deep Learning Research Platform to PyTorch

The company who developed Chainer changed to use PyTorch as their DL framework.

5. Appendix

5.1. Code

To avoid duplicate code, I isolated the model's code into independent file and import it in train and inference code.

- Code architecture

```
$ tree . -I 'env|result|__pycache__|images|*.md|*.pdf'
.
├── inference.py
├── resnet_cifar10.py
└── train.py
```

- Model (resnet_cifar10.py)
 - Reference:
 1. https://github.com/akamaster/pytorch_resnet_cifar10/blob/master/resnet.py
 2. <https://github.com/mitmul/chainer-cifar10/blob/master/models/ResNet.py>

```
#!/usr/bin/env python
# -*- coding: utf-8 -*-
import chainer
import chainer.functions as F
import chainer.links as L

class BottleNeck(chainer.Chain):
    def __init__(self, n_in, n_out, stride=1):
        self.shortcut = stride != 1
        super(BottleNeck, self).__init__()
        with self.init_scope():
            self.conv1 = L.Convolution2D(
                n_in, n_out, ksize=3, stride=stride, pad=1, nobias=True
```

```

    )
    self.bn1 = L.BatchNormalization(n_out)
    self.conv2 = L.Convolution2D(
        n_out, n_out, ksize=3, stride=1, pad=1, nobias=True
    )
    self.bn2 = L.BatchNormalization(n_out)
    self.conv3 = L.Convolution2D(
        n_in, n_out, ksize=1, stride=stride, pad=0, nobias=True
    )
    self.bn3 = L.BatchNormalization(n_out)

    def __call__(self, x):
        h = F.relu(self.bn1(self.conv1(x)))
        h = self.bn2(self.conv2(h))
        if self.shortcut:
            h += self.bn3(self.conv3(x))
        h = F.relu(h)
        return h

class Block(chainer.ChainList):
    def __init__(self, n_in, n_out, n_bottlenecks, stride):
        super(Block, self).__init__()
        self.in_planes = n_in
        strides = [stride] + [1] * (n_bottlenecks - 1)
        for stride in strides:
            self.add_link(BottleNeck(self.in_planes, n_out, stride))
            self.in_planes = n_out

    def __call__(self, x):
        for f in self:
            x = f(x)
        return x

class ResNet(chainer.Chain):
    def __init__(self, n_class=10, n_blocks=[3, 3, 3]):
        super(ResNet, self).__init__()

        with self.init_scope():
            self.conv1 = L.Convolution2D(None, 16, 3, 1, 1, nobias=True)
            self.bn2 = L.BatchNormalization(16)
            self.res3 = Block(16, 16, n_blocks[0], 1)
            self.res4 = Block(16, 32, n_blocks[1], 2)
            self.res5 = Block(32, 64, n_blocks[2], 2)
            self.fc7 = L.Linear(64, n_class)

    def __call__(self, x):
        h = F.relu(self.bn2(self.conv1(x)))
        h = self.res3(h)
        h = self.res4(h)
        h = self.res5(h)
        h = F.average_pooling_2d(h, h.shape[2:])
        h = self.fc7(h)
        return h

class ResNet20(ResNet):
    def __init__(self, n_class=10):
        super(ResNet20, self).__init__(n_class, [3, 3, 3])

```

- Train script (`train.py`)
 - Reference: "Homework 1 - building_model.pdf"

```

from __future__ import print_function
import chainer
from chainer import training
from chainer.training import extensions
from chainer import serializers
import chainer.links as L
import argparse
from resnet_cifar10 import ResNet20

parser = argparse.ArgumentParser(description="Chainer example: Cifar-10")
parser.add_argument(
    "--batchsize",

```

```

        "-b",
        type=int,
        default=128,
        help="Number of images in each mini-batch",
    )
    parser.add_argument(
        "--epoch",
        "-e",
        type=int,
        default=10,
        help="Number of sweeps over the dataset to train",
    )
    parser.add_argument(
        "-gpu", "-g", type=int, default=0, help="GPU ID (negative value indicates CPU)"
    ) # Set the initial matrix format(numpy/cupy)
    parser.add_argument(
        "-out", "-o", default="result/4", help="Directory to output the result"
    )
    parser.add_argument(
        "--resume", "-r", default="", help="Resume the training from snapshot"
    )
    parser.add_argument(
        "--unit", "-u", type=int, default=10, help="Number of output layer units"
    )

if __name__ == "__main__":
    args = parser.parse_args(["-g", "-1"]) # The negative device number means CPU.
    print("GPU: {}".format(args.gpu))
    print("# unit: {}".format(args.unit))
    print("# Minibatch-size: {}".format(args.batchsize))
    print("# epoch: {}".format(args.epoch))
    print("")

    model = ResNet20(args.unit)
    classifier_model = L.Classifier(model)
    # if args.gpu >= 0:
    #     chainer.cuda.get_device(args.gpu).use() # Make a specified GPU current
    #     classifier_model.to_gpu() # Copy the model to the GPU

    optimizer = (
        chainer.optimizers.Adam()
    ) # Adam is one of the Gradient descent optimizationalgorithms.
    optimizer.setup(classifier_model)

    train, test = chainer.datasets.get_cifar10()
    train_iter = chainer.iterators.SerialIterator(
        train, args.batchsize
    ) # Set Training data batchiterater
    test_iter = chainer.iterators.SerialIterator(
        test, args.batchsize, repeat=False, shuffle=False
    )
    # Forward the test data after each of training to calcuat the validation loss/arruracy.
    # updater and trainer
    updater = training.StandardUpdater(train_iter, optimizer, device=args.gpu)
    trainer = training.Trainer(updater, (args.epoch, "epoch"), out=args.out)
    # extension objects
    trainer.extend(extensions.Evaluator(test_iter, classifier_model, device=args.gpu))
    trainer.extend(extensions.dump_graph("main/loss"))
    trainer.extend(extensions.snapshot(), trigger=(1, "epoch"))
    trainer.extend(extensions.LogReport())
    trainer.extend(
        extensions.PrintReport(
            [
                "epoch",
                "main/loss",
                "validation/main/loss",
                "main/accuracy",
                "validation/main/accuracy",
                "elapsed_time",
            ]
        )
    )
    trainer.extend(extensions.ProgressBar())

    if args.resume:
        # Resume from a snapshot
        serializers.load_npz(args.resume, trainer)
    trainer.run()
    serializers.save_npz(
        "{}resnet20.model".format(args.out), model
    )

```

```
) # Save the model(all trainedweights)
```

- Inference (inference.py)
 - Reference: "Homework 1 - building_model.pdf"

```
from resnet_cifar10 import ResNet20
import argparse
import chainer
from chainer import serializers
import numpy as np

parser = argparse.ArgumentParser(description="Chainer example: Cifar-10")
parser.add_argument(
    "--out", "-o", default="result/4/resnet20.model", help="Directory to output the result"
)
parser.add_argument(
    "--unit", "-u", type=int, default=10, help="Number of output layer units"
)
# parser.add_argument('--gpu', '-g', type=int, default=0,
# help='GPU ID (negative value indicates CPU)')#Set the initial matrixformat(numpy/cupy)
if __name__ == "__main__":
    args = parser.parse_args() # The negative device number means CPU.
    # print('GPU: {}'.format(args.gpu))
    print("# unit: {}".format(args.unit))
    print("")

    # Load the dataset
    _, test = chainer.datasets.get_cifar10()
    # Load trained model
    model = ResNet20(args.unit)
    # if args.gpu >= 0:
    # chainer.cuda.get_device(args.gpu).use() # Make a specified GPU current
    # model.to_gpu() # Copy the model to the GPU
    # xp = np if args.gpu < 0 else chainer.cuda.cupy
    xp = np

    serializers.load_npz(args.out, model)

    x = chainer.Variable(xp.asarray([test[0][0]])) # test data
    t = chainer.Variable(xp.asarray([test[0][1]])) # labels
    y = model(x) # Inference result
    print("The test data label:", xp.asarray([test[0][1]]))
    print("result:", y)
    y_top5 = y.array[0].argsort()[-5:][::-1]
    print("result Top 1:", [y_top5[0]])
    print("result Top 5:", y_top5)
```