

# Deep Learning System and Parallel Computing HW1

---

- Author: 黃柏瑄 (P78081528)

## 1. Experiment description

---

- Framework: Chainer
- Dataset: Cifar10
- Model: ResNet20

## 2. Environment setting

---

- Hardware:
  - CPU: Intel(R) Core(TM) i7-10700K CPU @ 3.80GHz (`$ lscpu | grep 'Model name'`)
  - Memory: 32 GB
- Software:
  - OS: Ubuntu 18.04.5 LTS (`$ lsb_release -d`)
  - Python 3.6.9
  - Chainer 7.7.0

```
$ virtualenv --python python3.6 env
$ source env/bin/activate
$ pip install chainer==7.7.0
$ pip list
Package          Version
-----
chainer           7.7.0
filelock          3.0.12
numpy             1.19.2
pip               20.2.2
protobuf          3.13.0
setuptools        50.2.0
six               1.15.0
typing-extensions 3.7.4.3
wheel             0.35.1
```

## 3. Result

Code can be found at [Appendix](#) section.

### 3.1. Train

```
$ python train.py
```

```
[env] .../benlab/pro/ms/dpc2020/hw1/ [main ↑1 •1 010 x] [20:33]
$ python train.py
GPU: -1
# unit: 10
# Minibatch-size: 128
# epoch: 10

epoch      main/loss    validation/main/loss    main/accuracy    validation/main/accuracy    elapsed_time
1          1.49377      1.21208                 0.447011         0.563884                   1022.99
2          1.067       1.24985                 0.616828         0.561511                   2049.73
3          0.887509    0.890478                0.686759         0.691456                   3092.96
4          0.757574    0.960231                0.733156         0.665941                   4115.04
5          0.662063    0.770418                0.767104         0.728936                   5068.98
6          0.587038    0.83906                 0.795533         0.715783                   5988.97
7          0.531226    0.729016                0.815018         0.754252                   6926.88
8          0.486883    0.706037                0.829788         0.766515                   7837.67
9          0.434821    0.65807                 0.848326         0.778283                   8782.04
10         0.399522    0.700709                0.860394         0.773339                   9703.1

[env] .../benlab/pro/ms/dpc2020/hw1/ [main ↑1 •1 07 x] [23:16]
```

### 3.2. Inference

```
$ python inference.py
```

```
$ python inference.py
# unit: 10

The test data label: [3]
result: variable([[ -2.0492222 -3.4841754 -1.5581383 -1.2748083 -2.167873
                  -1.9728453 -2.962061 -2.8587072 -3.0967042 -2.4422274]])
result Top 1: [3]
result Top 5: [3 2 5 0 4]
```

## 4. Insight

### 4.1. CPU utilization

The CPU utilization is good during training phase.

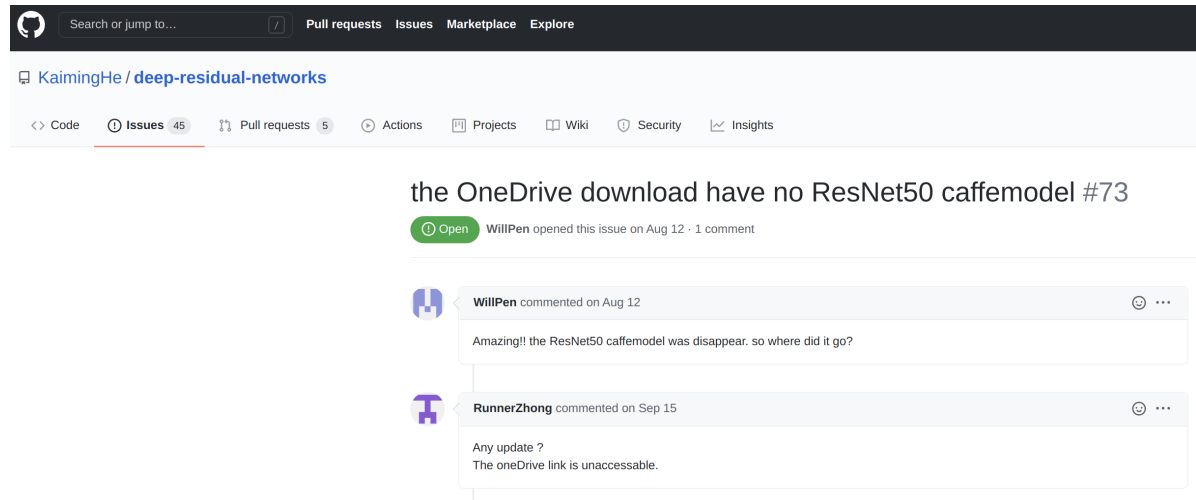
```
1  [|||||] 98.7% 5  [|||||] 98.7% 9  [|||||] 99.4% 13 [|||||] 98.7%
2  [|||||] 98.7% 6  [|||||] 99.4% 10 [|||||] 98.7% 14 [|||||] 99.3%
3  [|||||] 99.4% 7  [|||||] 98.7% 11 [|||||] 98.7% 15 [|||||] 98.7%
4  [|||||] 99.4% 8  [|||||] 98.7% 12 [|||||] 98.7% 16 [|||||] 100.0%
Mem[|||||] 18.2G/31.3G Tasks: 261, 1673 thr; 16 running
Swp[|||||] 754M/30.5G Load average: 11.39 8.58 4.37
Uptime: 19 days, 21:42:52

PID PRI VIRT RES SHR S CPU% MEM% TIME+ Command
14846 20 5949M 3905M 16680 R 1570 12.2 51:41.52 python train.py
15036 20 5949M 3905M 16680 R 99.1 12.2 2:57.07 python train.py
15032 20 5949M 3905M 16680 R 98.5 12.2 2:57.26 python train.py
15038 20 5949M 3905M 16680 R 98.5 12.2 2:55.99 python train.py
15028 20 5949M 3905M 16680 R 98.5 12.2 2:57.53 python train.py
15030 20 5949M 3905M 16680 R 98.5 12.2 2:57.35 python train.py
15025 20 5949M 3905M 16680 R 98.5 12.2 2:58.77 python train.py
15034 20 5949M 3905M 16680 R 98.5 12.2 2:56.84 python train.py
15033 20 5949M 3905M 16680 R 97.8 12.2 2:57.91 python train.py
15035 20 5949M 3905M 16680 R 97.8 12.2 2:56.92 python train.py
15027 20 5949M 3905M 16680 R 97.8 12.2 2:59.26 python train.py
15031 20 5949M 3905M 16680 R 97.8 12.2 2:56.96 python train.py
15037 20 5949M 3905M 16680 R 97.8 12.2 2:56.78 python train.py
15029 20 5949M 3905M 16680 R 97.2 12.2 2:57.57 python train.py
15039 20 5949M 3905M 16680 R 97.2 12.2 2:54.45 python train.py
15026 20 5949M 3905M 16680 R 95.9 12.2 2:58.77 python train.py
```

## 4.2. Fun facts

### 4.2.1. It is not easy to find some pre-trained weights on the Internet

Many pre-trained models are disappear. There may be many reasons such as privacy issue, business considerations, or lack of maintenance.



source: <https://github.com/KaimingHe/deep-residual-networks/issues/73>

### 4.2.2. Preferred Networks Migrates its Deep Learning Research Platform to PyTorch

The company who developed Chainer changed to use PyTorch as their DL framework.

## 5. Appendix

### 5.1. Code

To avoid duplicate code, I isolated the model's code into independent file and import it in train and inference code.

- Code architecture

```
$ tree . -I 'env|result|__pycache__|images|*.md|*.pdf'
.
├── inference.py
├── resnet_cifar10.py
└── train.py
```

- Model (resnet\_cifar10.py)
  - Reference:
    1. [https://github.com/akamaster/pytorch\\_resnet\\_cifar10/blob/master/resnet.py](https://github.com/akamaster/pytorch_resnet_cifar10/blob/master/resnet.py)
    2. <https://github.com/mitmul/chainer-cifar10/blob/master/models/ResNet.py>

```
#!/usr/bin/env python
# -*- coding: utf-8 -*-
import chainer
import chainer.functions as F
import chainer.links as L

class BottleNeck(chainer.Chain):
    def __init__(self, n_in, n_out, stride=1):
        self.shortcut = stride != 1
        super(BottleNeck, self).__init__()
        with self.init_scope():
            self.conv1 = L.Convolution2D(
                n_in, n_out, ksize=3, stride=stride, pad=1, nobias=True
```

```

    )
    self.bn1 = L.BatchNormalization(n_out)
    self.conv2 = L.Convolution2D(
        n_out, n_out, ksize=3, stride=1, pad=1, nobias=True
    )
    self.bn2 = L.BatchNormalization(n_out)
    self.conv3 = L.Convolution2D(
        n_in, n_out, ksize=1, stride=stride, pad=0, nobias=True
    )
    self.bn3 = L.BatchNormalization(n_out)

def __call__(self, x):
    h = F.relu(self.bn1(self.conv1(x)))
    h = self.bn2(self.conv2(h))
    if self.shortcut:
        h += self.bn3(self.conv3(x))
    h = F.relu(h)
    return h

class Block(chainer.ChainList):
    def __init__(self, n_in, n_out, n_bottlenecks, stride):
        super(Block, self).__init__()
        self.in_planes = n_in
        strides = [stride] + [1] * (n_bottlenecks - 1)
        for stride in strides:
            self.add_link(BottleNeck(self.in_planes, n_out, stride))
            self.in_planes = n_out

    def __call__(self, x):
        for f in self:
            x = f(x)
        return x

class ResNet(chainer.Chain):
    def __init__(self, n_class=10, n_blocks=[3, 3, 3]):
        super(ResNet, self).__init__()

        with self.init_scope():
            self.conv1 = L.Convolution2D(None, 16, 3, 1, 1, nobias=True)
            self.bn2 = L.BatchNormalization(16)
            self.res3 = Block(16, 16, n_blocks[0], 1)
            self.res4 = Block(16, 32, n_blocks[1], 2)
            self.res5 = Block(32, 64, n_blocks[2], 2)
            self.fc7 = L.Linear(64, n_class)

    def __call__(self, x):
        h = F.relu(self.bn2(self.conv1(x)))
        h = self.res3(h)
        h = self.res4(h)
        h = self.res5(h)
        h = F.average_pooling_2d(h, h.shape[2:])
        h = self.fc7(h)
        return h

class ResNet20(ResNet):
    def __init__(self, n_class=10):
        super(ResNet20, self).__init__(n_class, [3, 3, 3])

```

- Train script ( `train.py` )
  - Reference: "Homework 1 - building\_model.pdf"

```

from __future__ import print_function
import chainer
from chainer import training
from chainer.training import extensions
from chainer import serializers
import chainer.links as L
import argparse
from resnet_cifar10 import ResNet20

parser = argparse.ArgumentParser(description="Chainer example: Cifar-10")
parser.add_argument(
    "--batchsize",

```

```

        "-b",
        type=int,
        default=128,
        help="Number of images in each mini-batch",
    )
    parser.add_argument(
        "--epoch",
        "-e",
        type=int,
        default=10,
        help="Number of sweeps over the dataset to train",
    )
    parser.add_argument(
        "-gpu", "-g", type=int, default=0, help="GPU ID (negative value indicates CPU)"
    ) # Set the initial matrixformat(numpy/cupy)
    parser.add_argument(
        "-out", "-o", default="result/4", help="Directory to output the result"
    )
    parser.add_argument(
        "--resume", "-r", default="", help="Resume the training from snapshot"
    )
    parser.add_argument(
        "--unit", "-u", type=int, default=10, help="Number of output layer units"
    )

if __name__ == "__main__":
    args = parser.parse_args(["-g", "-1"]) # The negative device number means CPU.
    print("GPU: {}".format(args.gpu))
    print("# unit: {}".format(args.unit))
    print("# Minibatch-size: {}".format(args.batchsize))
    print("# epoch: {}".format(args.epoch))
    print("")

    model = ResNet20(args.unit)
    classifier_model = L.Classifier(model)
    # if args.gpu >= 0:
    # chainer.cuda.get_device(args.gpu).use() # Make a specified GPU current
    # classifier_model.to_gpu() # Copy the model to the GPU

    optimizer = (
        chainer.optimizers.Adam()
    ) # Adam is one of the Gradient descent optimizationalgorithms.
    optimizer.setup(classifier_model)

    train, test = chainer.datasets.get_cifar10()
    train_iter = chainer.iterators.SerialIterator(
        train, args.batchsize
    ) # Set Training data batchiterater
    test_iter = chainer.iterators.SerialIterator(
        test, args.batchsize, repeat=False, shuffle=False
    )
    # Forward the test data after each of training to calcuat the validation loss/arruracy.
    # updater and trainer
    updater = training.StandardUpdater(train_iter, optimizer, device=args.gpu)
    trainer = training.Trainer(updater, (args.epoch, "epoch"), out=args.out)
    # extension objects
    trainer.extend(extensions.Evaluator(test_iter, classifier_model, device=args.gpu))
    trainer.extend(extensions.dump_graph("main/loss"))
    trainer.extend(extensions.snapshot(), trigger=(1, "epoch"))
    trainer.extend(extensions.LogReport())
    trainer.extend(
        extensions.PrintReport(
            [
                "epoch",
                "main/loss",
                "validation/main/loss",
                "main/accuracy",
                "validation/main/accuracy",
                "elapsed_time",
            ]
        )
    )
    trainer.extend(extensions.ProgressBar())

    if args.resume:
        # Resume from a snapshot
        serializers.load_npz(args.resume, trainer)
    trainer.run()
    serializers.save_npz(
        "{}resnet20.model".format(args.out), model
    )

```

```
) # Save the model(all trainedweights)
```

- Inference (inference.py)
  - Reference: "Homework 1 - building\_model.pdf"

```
from resnet_cifar10 import ResNet20
import argparse
import chainer
from chainer import serializers
import numpy as np

parser = argparse.ArgumentParser(description="Chainer example: Cifar-10")
parser.add_argument(
    "--out", "-o", default="result/4/resnet20.model", help="Directory to output the result"
)
parser.add_argument(
    "--unit", "-u", type=int, default=10, help="Number of output layer units"
)
# parser.add_argument('--gpu', '-g', type=int, default=0,
# help='GPU ID (negative value indicates CPU)')#Set the initial matrixformat(numpy/cupy)
if __name__ == "__main__":
    args = parser.parse_args() # The negative device number means CPU.
    # print('GPU: {}'.format(args.gpu))
    print("# unit: {}".format(args.unit))
    print("")

    # Load the dataset
    _, test = chainer.datasets.get_cifar10()
    # Load trained model
    model = ResNet20(args.unit)
    # if args.gpu >= 0:
    # chainer.cuda.get_device(args.gpu).use() # Make a specified GPU current
    # model.to_gpu() # Copy the model to the GPU
    # xp = np if args.gpu < 0 else chainer.cuda.cupy
    xp = np

    serializers.load_npz(args.out, model)

    x = chainer.Variable(xp.asarray([test[0][0]])) # test data
    t = chainer.Variable(xp.asarray([test[0][1]])) # labels
    y = model(x) # Inference result
    print("The test data label:", xp.asarray([test[0][1]]))
    print("result:", y)
    y_top5 = y.array[0].argsort()[-5:][::-1]
    print("result Top 1:", [y_top5[0]])
    print("result Top 5:", y_top5)
```